


# Thinking in Frequency


Slides adapted from James Hays, Derek Hoiem, Alyosha Efros, and Steven  
Lehar

**Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?**


Gaussian



Box filter



# Why does a lower resolution image still make sense to us? What do we lose?



# How is it that a 4MP image can be compressed to a few hundred KB without a noticeable change?



RGB image:

- 1 byte / sub-pixel
- \* 3 sub-pixels / pixel
- \* 4M pixels

---

12M bytes


# Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

*Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.*


*...the manner in which the author arrives at these equations is not exempt of difficulties and...his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.*

- Don't believe it?
  - Neither did Lagrange, Laplace, Poisson and other big wigs
  - Not translated into English until 1878!
- But it's (mostly) true!
  - called Fourier Series
  - there are some subtle restrictions



# Frequency Spectra


- example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$




Our building block:

$$A \sin(\omega x + \phi)$$


Add enough of them to get any signal  $f(x)$  you want!



# Frequency Spectra

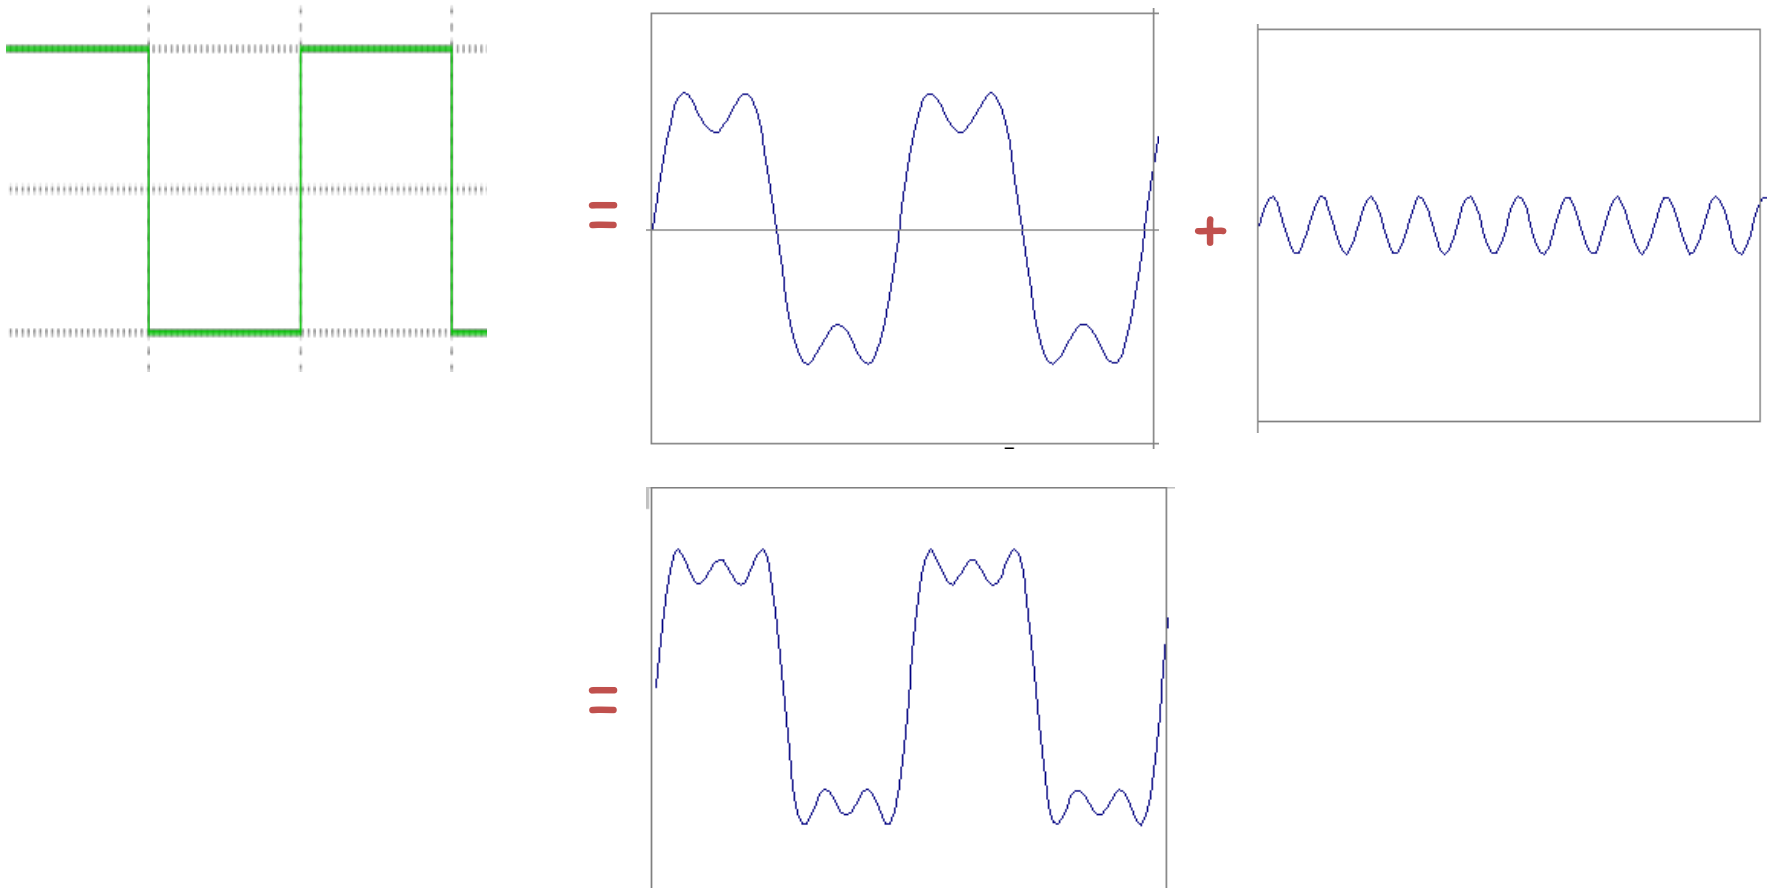


# Frequency Spectra

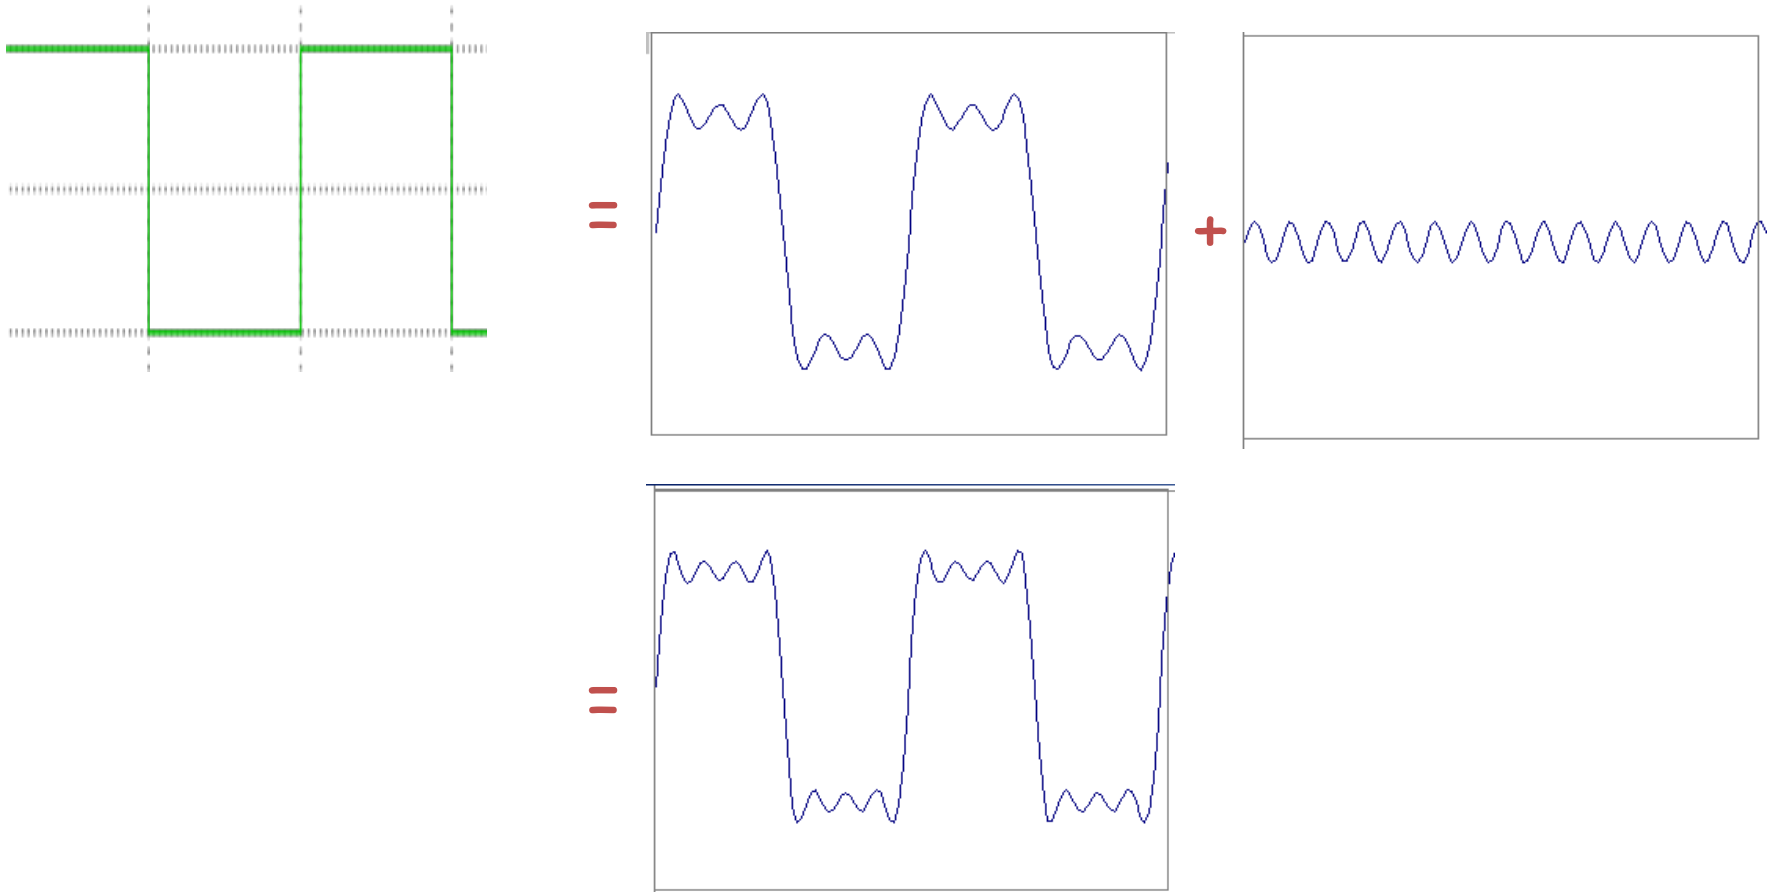





# Frequency Spectra




# Frequency Spectra




# Frequency Spectra




# Frequency Spectra






# Frequency Spectra



$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$











# Fourier analysis in images

- We can also think of all kinds of other signals the same way


Intensity Image



Fourier Image






# Frequency <-> Spatial







# Frequency $\leftrightarrow$ Spatial

**Brightness Image    Fourier Transform    Inverse Transformed**





**Low-Pass Filtered    Inverse Transformed**




# Frequency <-> Spatial


**High-Pass Filtered Inverse Transformed**



**Band-Pass Filtered Inverse Transformed**



# Signals can be composed



# The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$F[g * h] = F[g]F[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

$$F^{-1}[gh] = F^{-1}[g] * F^{-1}[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!


# Properties of Fourier Transforms

- Linearity  $\mathcal{F}[ax(t) + by(t)] = a\mathcal{F}[x(t)] + b\mathcal{F}[y(t)]$
- Fourier transform of a real signal is symmetric about the origin
- Can be computed efficiently
  - Fast Fourier Transform (FFT):  $O(N \log N)$

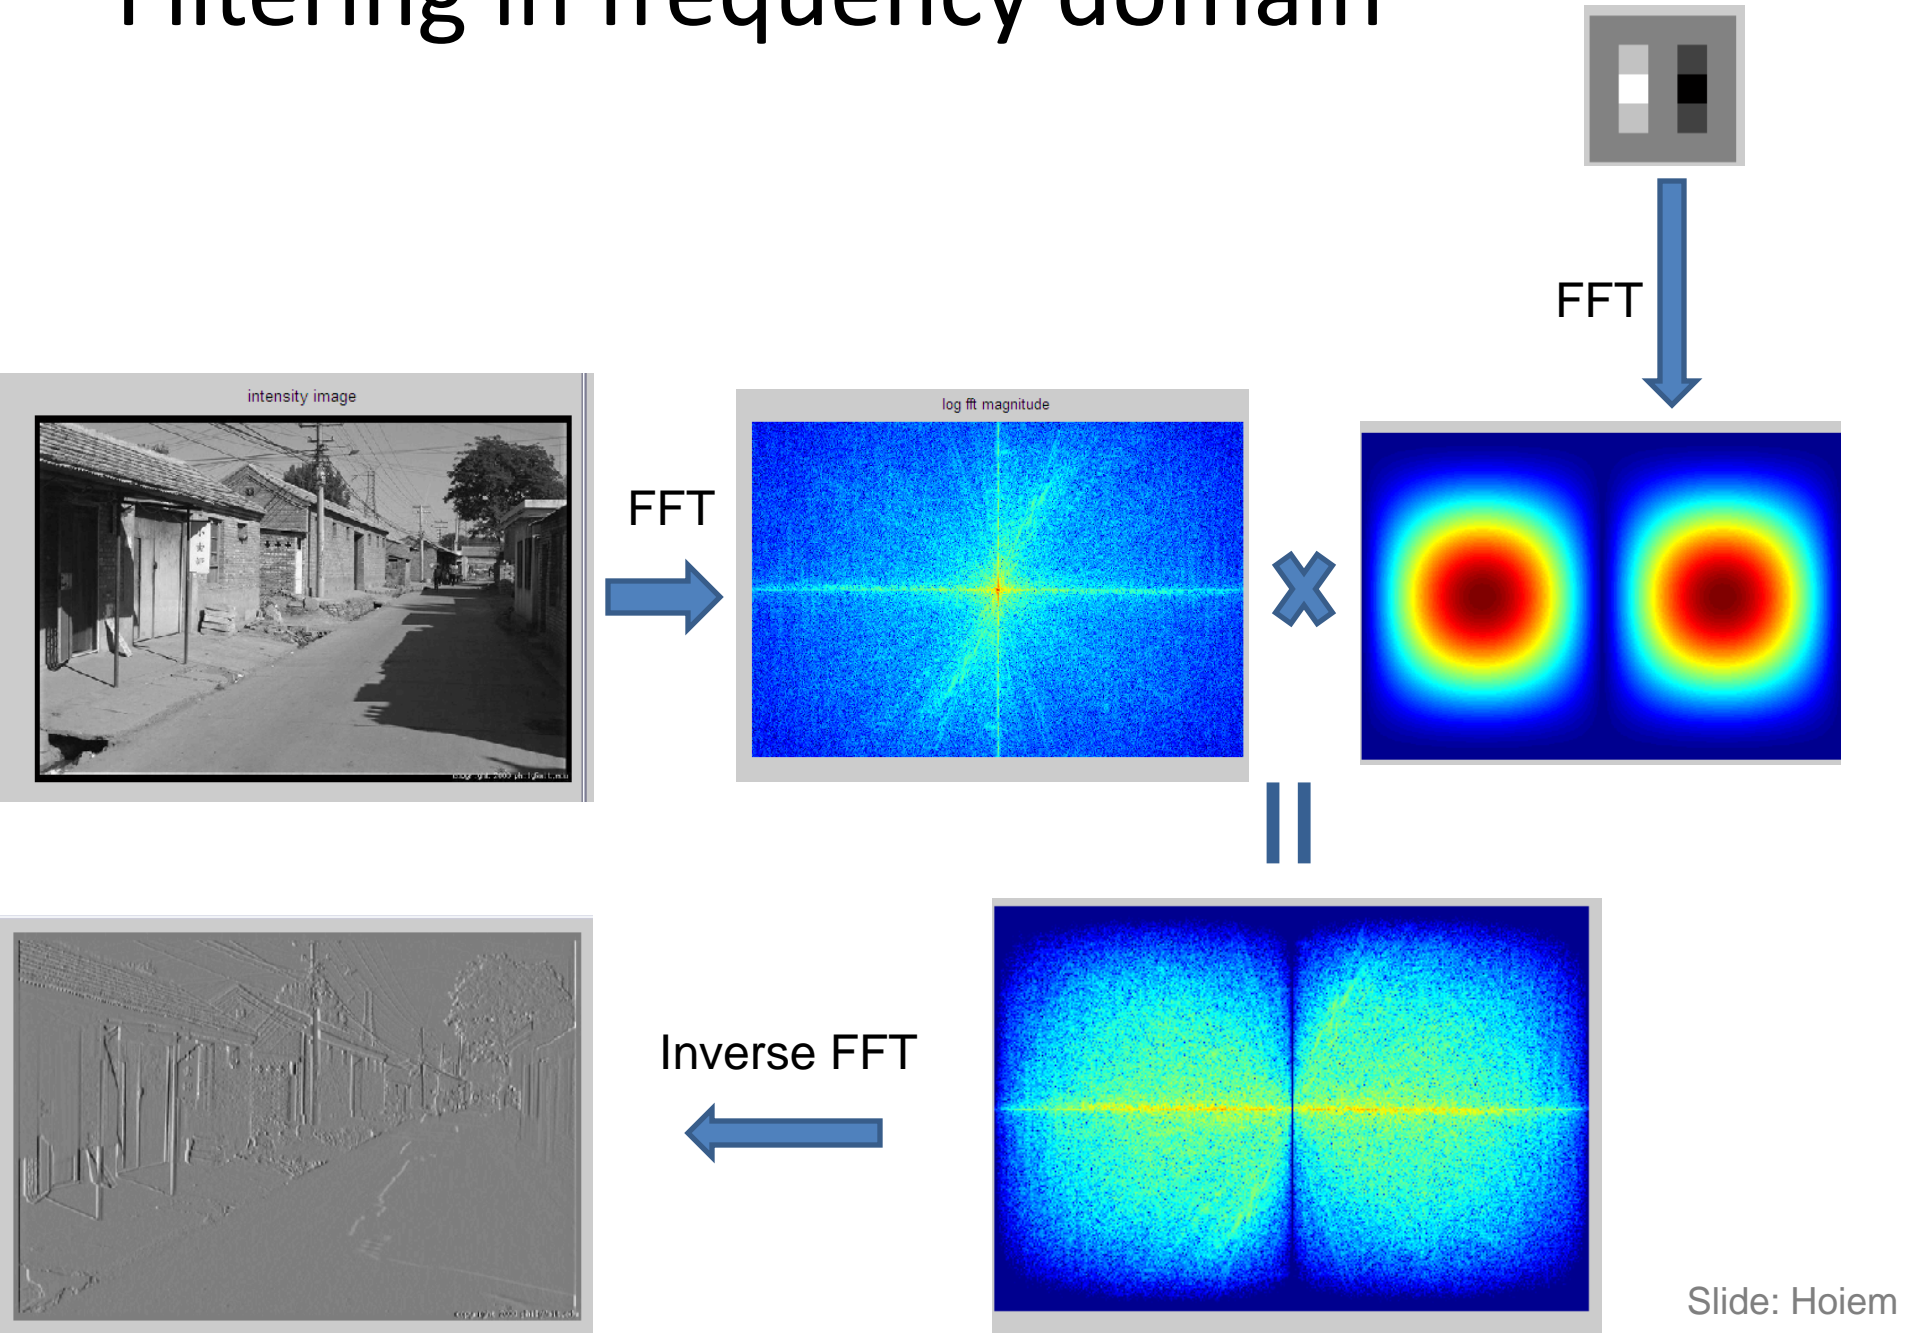
# Filtering in spatial domain

|   |   |    |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

intensity image



# Filtering in frequency domain



# FFT in Matlab

- Filtering with fft

```
im = double(imread('...'))/255;
im = rgb2gray(im); % "im" should be a gray-scale floating point image
[imh, imw] = size(im);

hs = 50; % filter half-size
fil = fspecial('gaussian', hs*2+1, 10);

fftsize = 1024; % should be order of 2 (for speed) and include padding
im_fft = fft2(im, fftsize, fftsize); % 1) fft im with padding
fil_fft = fft2(fil, fftsize, fftsize); % 2) fft fil, pad to same size as
image
im_fil_fft = im_fft .* fil_fft; % 3) multiply fft images
im_fil = ifft2(im_fil_fft); % 4) inverse fft2
im_fil = im_fil(1+hs:size(im,1)+hs, 1+hs:size(im, 2)+hs); % 5) remove padding
```

- Displaying with fft


```
figure(1), imagesc(log(abs(fftshift(im_fft)))), axis image, colormap jet
```




# Filtering

**Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?**

Gaussian



Box filter




# Gaussian


intensity image




filter: gaussian




filtered image




log fft magnitude of image



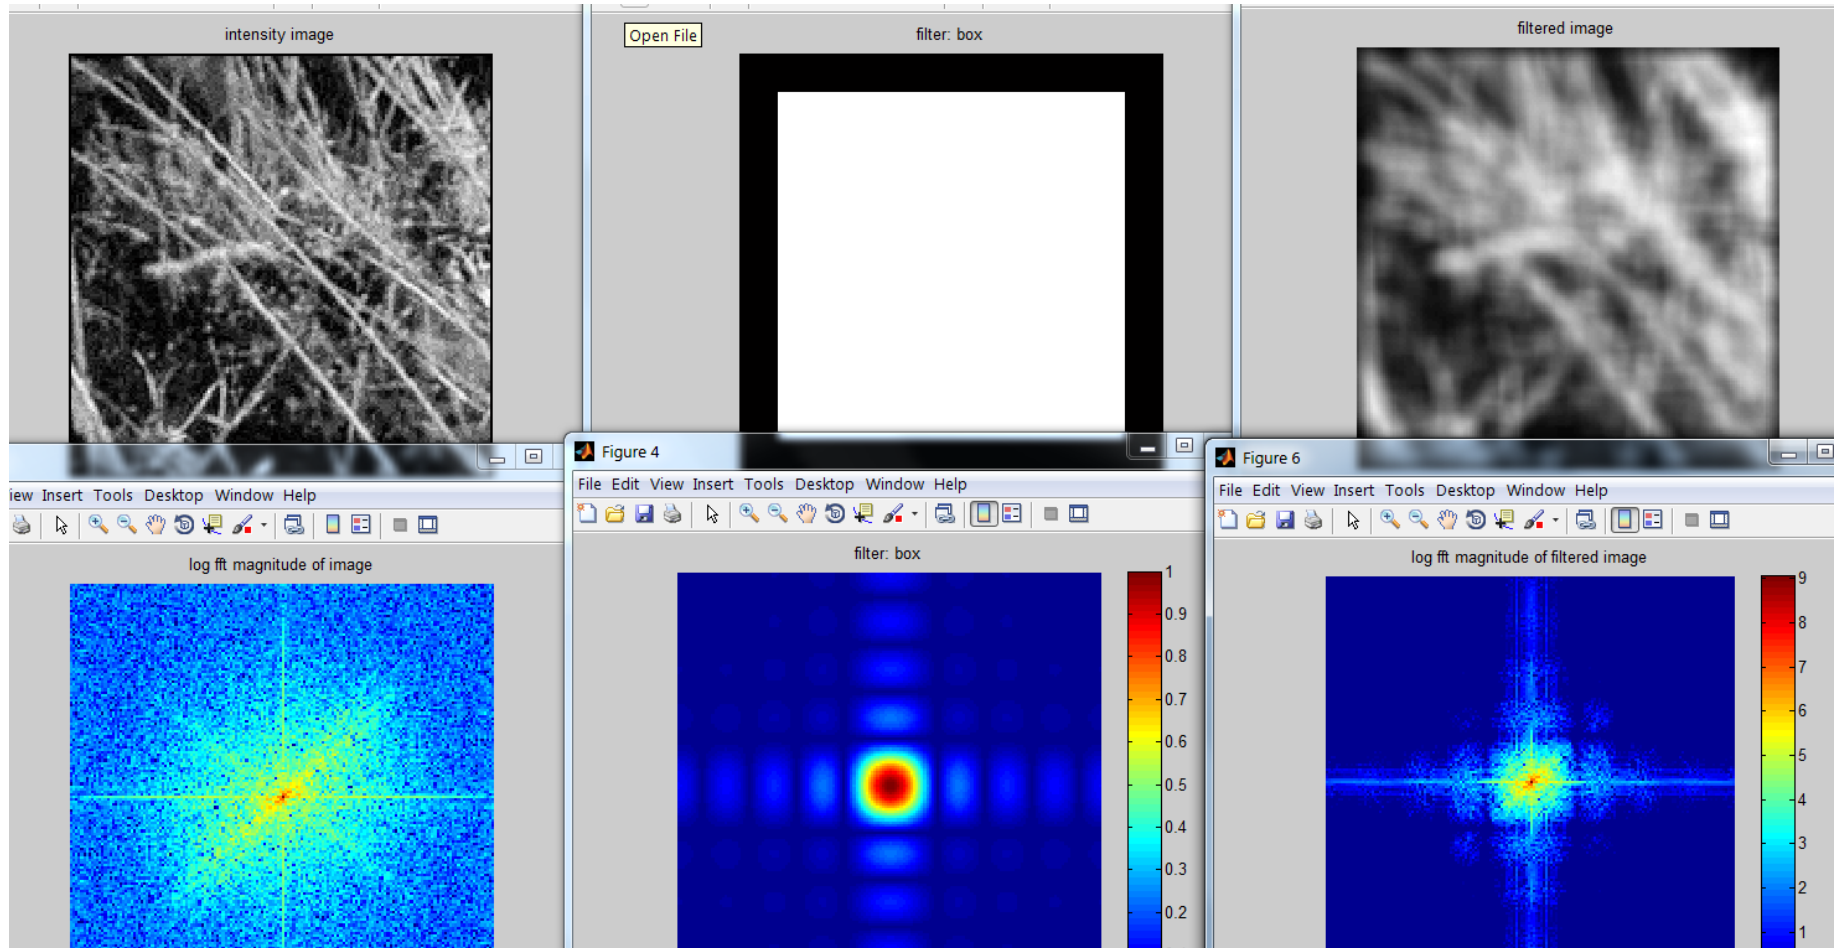
filter: gaussian



log fft magnitude of filtered image




# Box Filter



# Sampling

**Why does a lower resolution image still make sense to us? What do we lose?**




# Subsampling by a factor of 2



Throw away every other row and column to create a 1/2 size image


# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

- Sub-sampling may be dangerous....
- Characteristic errors may appear:
  - “Wagon wheels rolling the wrong way in movies”
  - “Checkerboards disintegrate in ray tracing”




# Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time =  $1/30$  sec. for video,  $1/24$  sec. for film):




Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)

# Aliasing in graphics



# Nyquist-Shannon Sampling Theorem

- When sampling a signal at discrete intervals, the sampling frequency must be  $\geq 2 \times f_{\max}$
- $f_{\max}$  = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



# Anti-aliasing


## Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
  - Will lose information
  - But it's better than aliasing
  - Apply a smoothing filter

# Algorithm for downsampling by factor of 2

1. Start with image(h, w)
2. Apply low-pass filter  
`im_blur = imfilter(image, fspecial('gaussian', 7, 1))`
3. Sample every other pixel  
`im_small = im_blur(1:2:end, 1:2:end);`

# Subsampling without pre-filtering



$1/2$



$1/4$  (2x zoom)




$1/8$  (4x zoom)

# Subsampling with Gaussian pre-filtering



Gaussian  $1/2$







G  $1/4$



G  $1/8$

# Things to Remember

- Sometimes it makes sense to think of images and filtering in the frequency domain
  - Fourier analysis
- Can be faster to filter using FFT for large images ( $N \log N$  vs.  $N^2$  for auto-correlation)
- Images are mostly smooth
  - Basis for compression
- Remember to low-pass before sampling

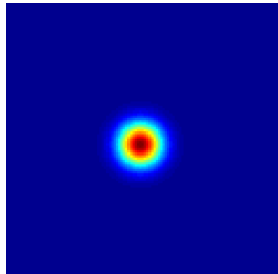




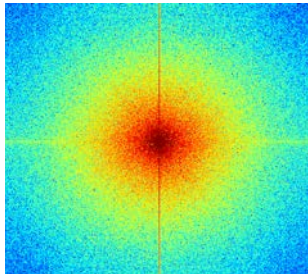
# Practice question

1. Match the spatial domain image to the Fourier magnitude image

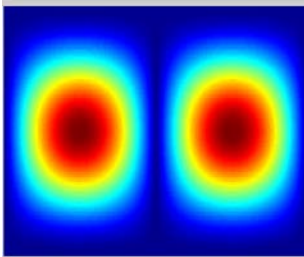
1



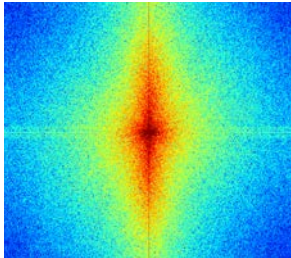
2



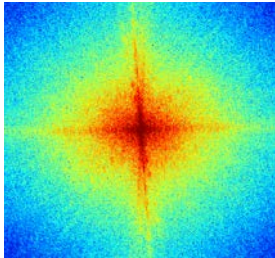
3




4




5



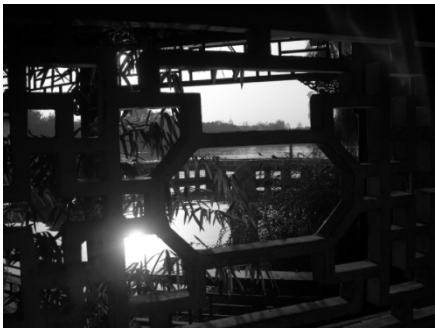
A



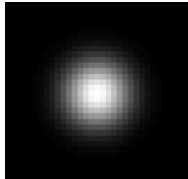
B




C



D



E



# Next class

- Template matching
- Image Pyramids
- Filter banks and texture
- Denoising, Compression